

## ▼ MINIOR PROJECT

### Problem Statement :

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. Several constraints were placed on the selection of these instances from a larger database. In particular, all patients here are females at least 21 years old of Pima Indian heritage.

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from sklearn.metrics import roc_curve, auc
6 from imblearn.over_sampling import SMOTE
7 from sklearn.model_selection import train_test_split
8 from sklearn import preprocessing
9
10 from sklearn.model_selection import GridSearchCV
11 from sklearn.model_selection import RandomizedSearchCV
12 from sklearn import metrics
13 from sklearn.metrics import roc_curve, auc
14
15 from sklearn.linear_model import LogisticRegression
16 from sklearn.naive_bayes import GaussianNB
17 from sklearn.neighbors import KNeighborsClassifier
18 from sklearn.svm import SVC
19
20
21 df = pd.read_excel("diabetes.csv.xlsx")
```

### Exploratory Data Analysis

```
1 # DISPLAY FIRST FIVE RECORDS OF DATA
2 df.head(5)
```

```
1 # Display last five records of data
2 df.tail(5)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedig
<b>763</b>	10	101	76	48	180	32.9	
<b>764</b>	2	122	70	27	0	36.8	
<b>765</b>	5	121	72	23	112	26.2	
<b>766</b>	1	126	60	0	0	30.1	
<b>767</b>	1	93	70	31	0	30.4	

```
1 # Display randomlly any number of records of data
2 df.sample(5)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedig
<b>15</b>	7	100	0	0	0	30.0	
<b>330</b>	8	118	72	19	0	23.1	
<b>461</b>	1	71	62	0	0	21.8	
<b>102</b>	0	125	96	0	0	22.5	
<b>352</b>	3	61	82	28	0	34.4	

## The shape of the dataset

```
1 #number of rows and columns
2 df.shape
```

```
(768, 9)
```

No. of Rows = 768

No. of Columns = 9

## List of all Columns

```
1 #list the types of all columns
2 df.dtypes
```

```
Pregnancies      int64
Glucose           int64
BloodPressure     int64
SkinThickness     int64
Insulin           int64
BMI               float64
```

```
DiabetesPedigreeFunction    float64
Age                        int64
Outcome                    int64
dtype: object
```

## Info of the dataset

```
1 #finding out if the dataset contains any null values
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
4   Insulin                768 non-null    int64
5   BMI                   768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                   768 non-null    int64
8   Outcome                768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

## Summary of the dataset

```
1 # Statistical summary
2 df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000

## Observation :

In the above table,the main value of columns

'Glucose','BloodPressure','skinthickness','insuline','BMI'is zero(0).It is clear that this values cant be

zero. So i am going to impute mean values of these respective columns insted of zero.

## Data cleaning

```
1 #check the shape before drop the duplicates
2 df.shape
```

```
(768, 9)
```

```
1 df=df.drop_duplicates()
```

```
1 # check the shape after drop the duplicates
2 df.shape
```

```
(768, 9)
```

Before drop and after drop the duplicates the data set has same shape which means no duplicates in the dataset.

## Check the Null Values

```
1 #count of null values
2 #check the missing values in any column
3 # #display number of null values is very column in dataset
4 df.isnull().sum()
```

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age              0
Outcome           0
dtype: int64
```

There is no null values in the given dataset.

```
1 df.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

## Check the no. of zero values in dataset.

```
1 print('No. of zero values in Glucose',df[df['Glucose']==0].shape[0])
```

No. of zero values in Glucose 5

```
1 print('no. of zero values in Bloodpressure',df[df['BloodPressure']==0].shape[0])
```

no. of zero values in Bloodpressure 35

```
1 print('no. of zero values in skinThickness',df[df['SkinThickness']==0].shape[0])
```

no. of zero values in skinThickness 227

```
1 print('no. of zero values in Insulin',df[df['Insulin']==0].shape[0])
```

no. of zero values in Insulin 374

```
1 print('no. of zero values in BMI',df[df['BMI']==0].shape[0])
```

no. of zero values in BMI 11

## Replace no. of zero values with mean of the columns

```
1 df['Glucose']=df['Glucose'].replace(0,df['Glucose'].mean())
2 print('No.of zero values in Glucose',df[df['Glucose']==0].shape[0])
```

No.of zero values in Glucose 0

```
1 df['BloodPressure']=df['BloodPressure'].replace(0,df['BloodPressure'].mean())
2 df['SkinThickness']=df['SkinThickness'].replace(0,df['SkinThickness'].mean())
3 df['Insulin']=df['Insulin'].replace(0,df['Insulin'].mean())
4 df['BMI']=df['BMI'].replace(0,df['BMI'].mean())
```

```
1 df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	121.681605	72.254807	26.606479	118.660163	32.450805
<b>std</b>	3.369578	30.436016	12.115932	9.631241	93.080358	6.875374
<b>min</b>	0.000000	44.000000	24.000000	7.000000	14.000000	18.200000
<b>25%</b>	1.000000	99.750000	64.000000	20.536458	79.799479	27.500000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	79.799479	32.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000

## Count plot

```



1 #outcome count plot
2 f,ax=plt.subplots(1,2,figsize=(10,5))
3 df['Outcome'].value_counts().plot.pie(explode=[0,0.1],autopct='%1.1f%%',ax=ax[0],shadow
4 ax[0].set_title('Outcome')
5 ax[0].set_ylabel('')
6 sns.countplot('Outcome',data=df,ax=ax[1])
7 ax[1].set_title('Outcome')
8 N,P = df['Outcome'].value_counts()
9 print('Negative (0): ',N)
10 print('Positive (1): ',P)
11 plt.grid()
12 plt.show()

```

```

Negative (0):  500
Positive (1):  268

```



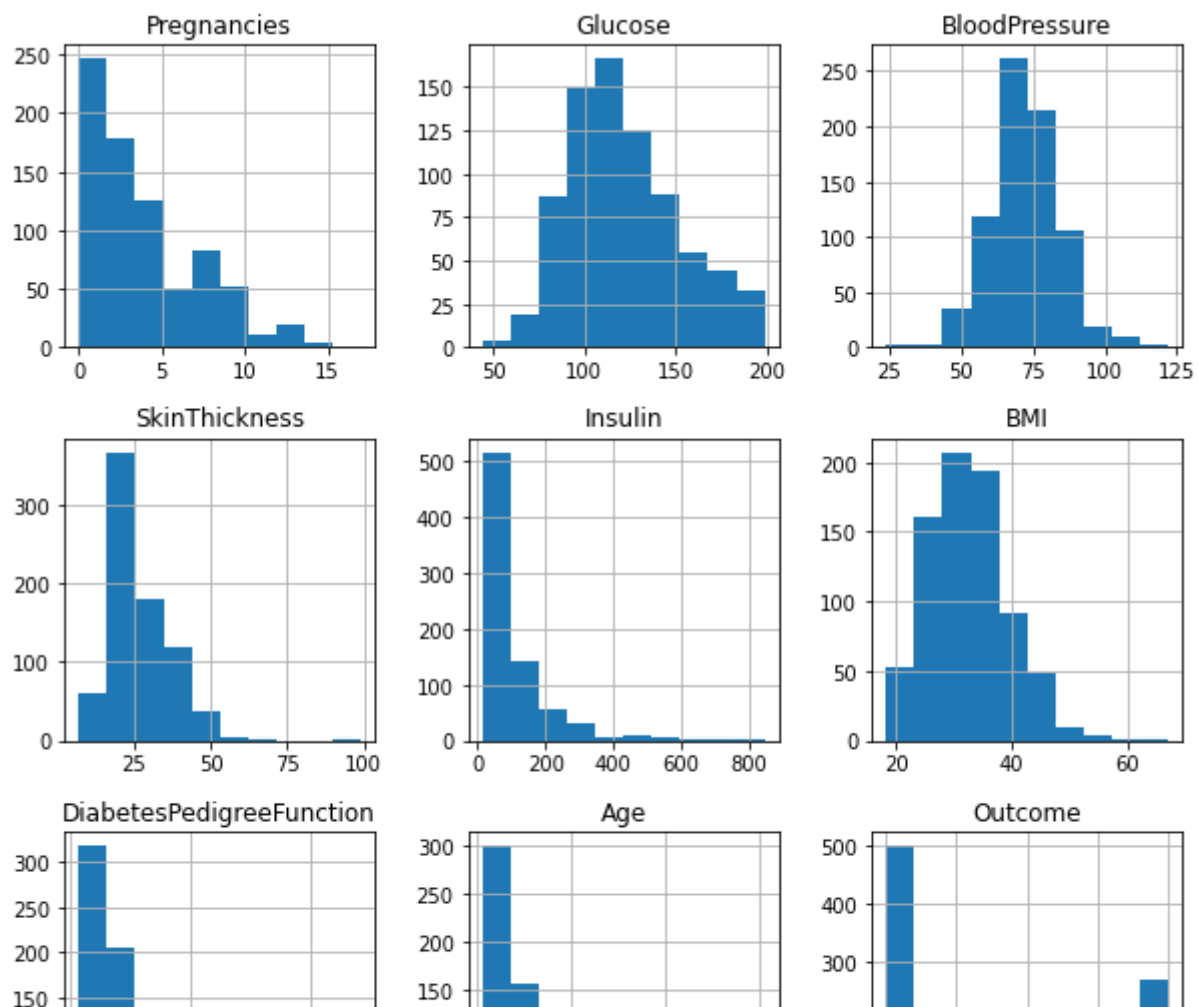
Out of total 768 people, 268 are diabetic (Positive(1)) and 500 are non-diabetic (Negative(0)). In the outcome column, **1** represents **diabetes positive** and **0** represents **diabetes negative**. The countplot tells us that the dataset is Imbalanced, as number of patients who don't have diabetes is more than those who have diabetes.

## Histograms

```

1 #Histogram of each feature
2 df.hist(bins=10,figsize=(10,10))
3 plt.show()

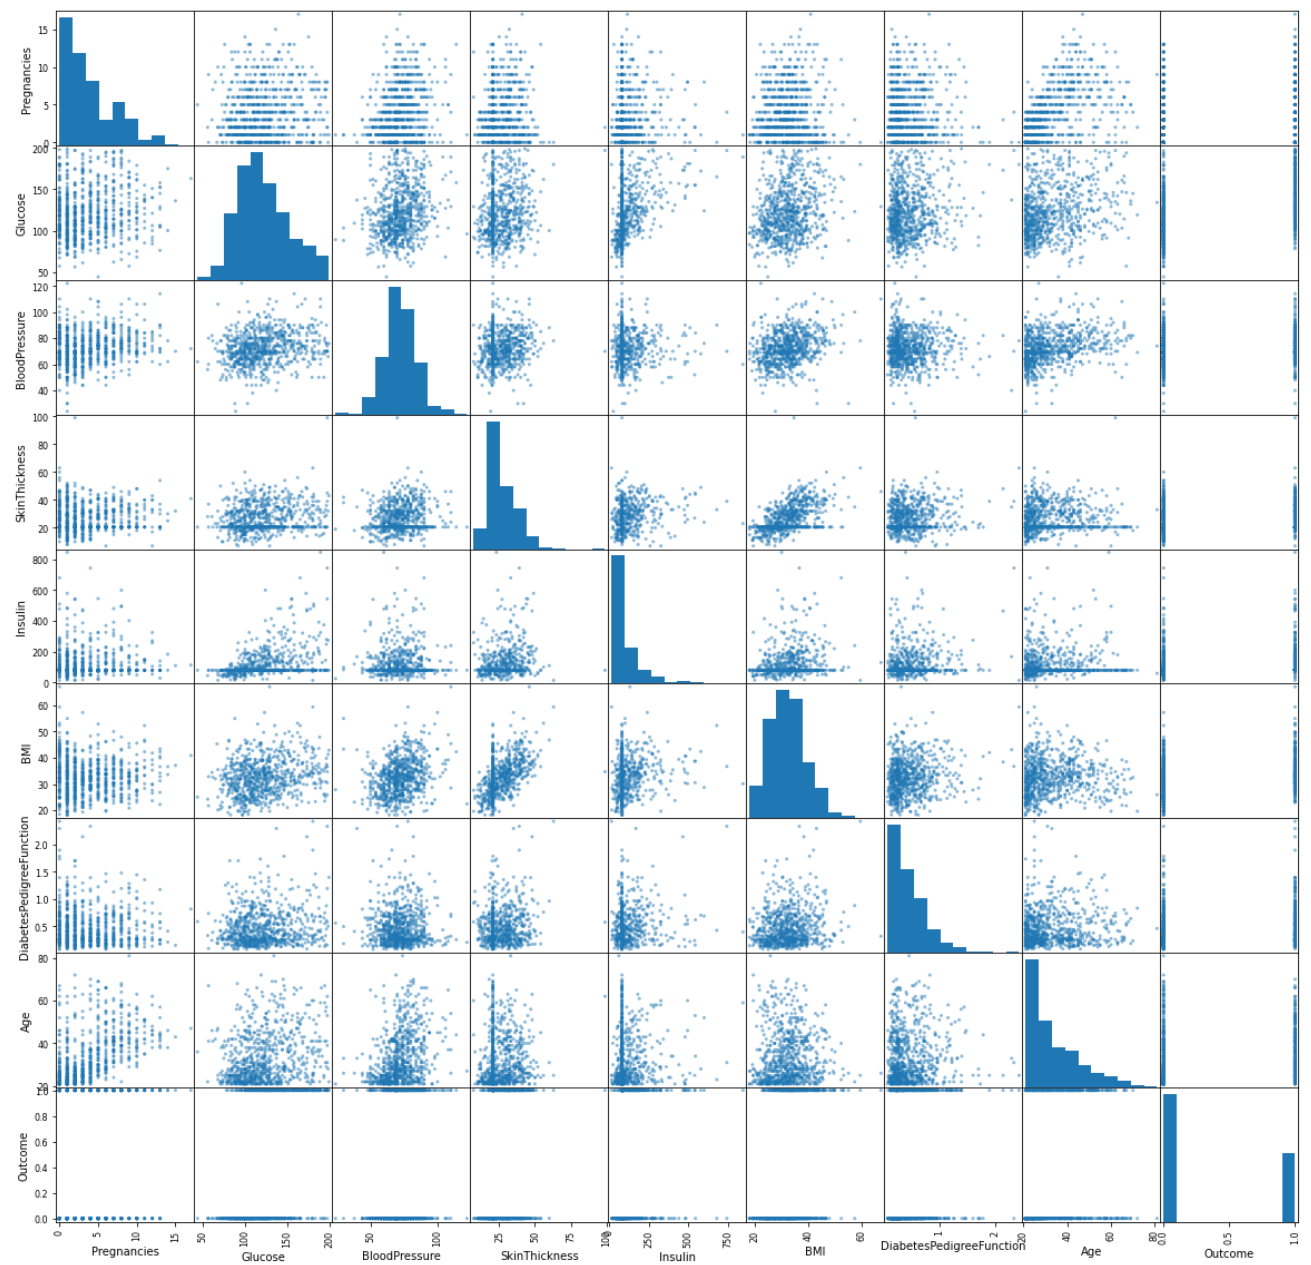
```



## Scatter plot



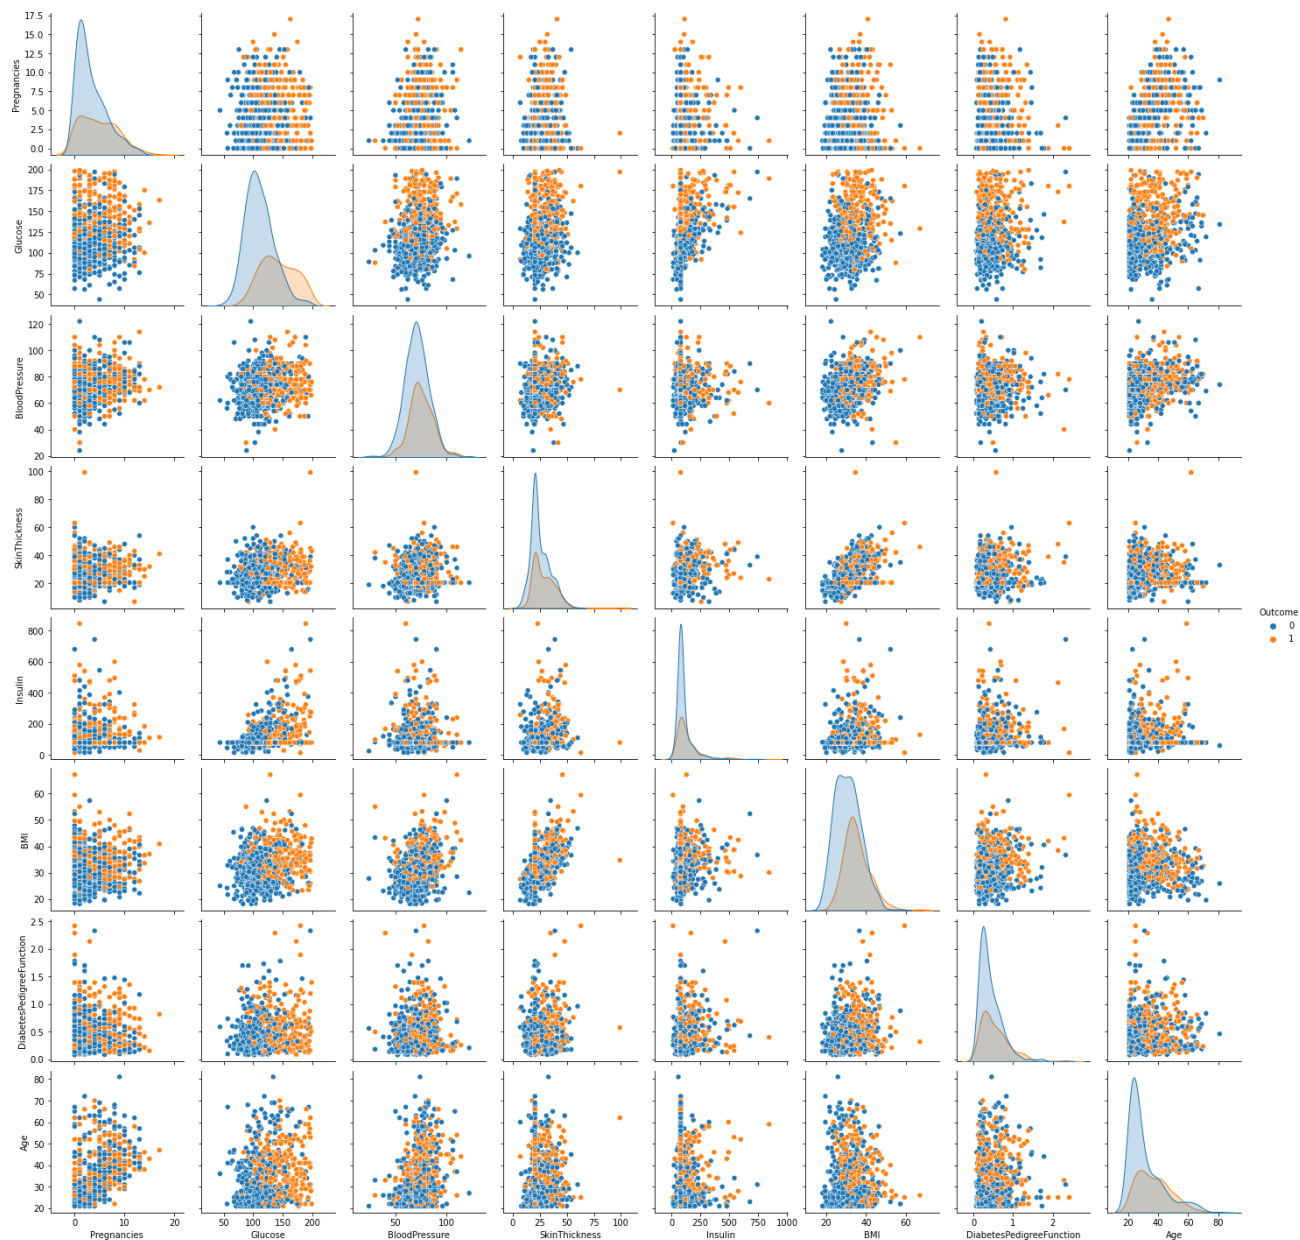
```
1 # Scatter plot matrix
2 from pandas.plotting import scatter_matrix
3 scatter_matrix(df,figsize=(20,20));
```



## Pair plot

```
1 #Pairplot
2 sns.pairplot(data = df, hue = 'Outcome')
3 plt.show()
```






Analyzing relationships between variables

Correlation analysis

```

1 import seaborn as sns
2 #get correlation of each features in dataset
3 corrmatrix = df.corr()
4 top_corr_features = corrmatrix.index
5 plt.figure(figsize = (10,10))
6 #plot heat map
7 g=sns.heatmap(df[top_corr_features].corr(),annot=True,cmap="RdYlGn")

```



### Observations:

From the correlation heatmap, we can see that there is a high correlation between Outcome and [Pregnancies, Glucose, BMI, Age, Insulin]. We can select these features to accept input from the user and predict the outcome.

## Split the data frame into X & y

```
1 target_name = 'Outcome'
2
3 # Separate object for target feature
4 y = df[target_name]
5
6 # Separate Object for Input Features
7 X = df.drop(target_name, axis=1)
```

```
1 X.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedi
0	6	148.0	72.0	35.000000	79.799479	33.6	
1	1	85.0	66.0	29.000000	79.799479	26.6	
2	8	183.0	64.0	20.536458	79.799479	23.3	
3	1	89.0	66.0	23.000000	94.000000	28.1	
4	0	137.0	40.0	35.000000	168.000000	43.1	

```
1 y.head()
```

```
0    1
1    0
2    1
3    0
4    1
Name: Outcome, dtype: int64
```

## Apply Feature Scaling

```
1 #Apply Standard Scaler
2 from sklearn.preprocessing import StandardScaler
3 scaler = StandardScaler()
4 scaler.fit(X)
5 SSX = scaler.transform(X)
```

## Train Test Split

```
1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(SSX, y , test_size=0.2, random_stat
```

```
1 X_train.shape,y_train.shape
```

```
((614, 8), (614,))
```

```
1 X_test.shape,v test.shape
```

```
((154, 8), (154,))
```

## Bulid the Classification Algorithms

### Logistic Regression

```
1 from sklearn.linear_model import LogisticRegression
2 lr = LogisticRegression(solver='liblinear',multi_class='ovr')
3 lr.fit(X_train,y_train)

LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)
```

### KNeighbors(KNN)

```
1 from sklearn.neighbors import KNeighborsClassifier
2 knn=KNeighborsClassifier()
3 knn.fit(X_train,y_train)

KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                    weights='uniform')
```

### \*Naive-Bayes \*

```
1 from sklearn.naive_bayes import GaussianNB
2 nb=GaussianNB()
3 nb.fit(X_train,y_train)

GaussianNB(priors=None, var_smoothing=1e-09)
```

### Support Vector Machine(SVM)

```
1 from sklearn.svm import SVC
2 sv=SVC()
3 sv.fit(X_train,y_train)

SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

### \*Decision Tree \*

```
1 from sklearn.tree import DecisionTreeClassifier
2 dt=DecisionTreeClassifier()
3 dt.fit(X_train,y_train)

DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                      max_depth=None, max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort='deprecated',
                      random_state=None, splitter='best')
```

## Random Forest

```
1 from sklearn.ensemble import RandomForestClassifier
2 rf=RandomForestClassifier(criterion='entropy')
3 rf.fit(X_train,y_train)

RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                      criterion='entropy', max_depth=None, max_features='auto',
                      max_leaf_nodes=None, max_samples=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=100,
                      n_jobs=None, oob_score=False, random_state=None,
                      verbose=0, warm_start=False)
```

## Making Prediction

```
1 X_test.shape

(154, 8)
```

## Making Prediction using Logistic Regression

```
1 ## Making predictions on test dataset
2 lr_pred=lr.predict(X_test)

1 lr_pred.shape

(154,)
```

## Making Prediction using KNN

```
1 ## Making predictions on test dataset
2 knn_pred=knn.predict(X_test)

1 knn_pred.shape

(154,)
```

## Making Prediction using Naive-Byes

```
1 ## Making predictions on test dataset
2 nb_pred=nb.predict(X_test)
```

```
1 nb_pred.shape

(154,)
```

## Making Prediction using SVM

```
1 #Making predictions on test dataset
2 sv_pred=sv.predict(X_test)
```

```
1 sv_pred.shape

(154,)
```

## Making Prediction using Decision Tree

```
1 #Making predictions on test dataset
2 dt_pred=dt.predict(X_test)
```

```
1 dt_pred.shape

(154,)
```

## Making Prediction using Random Forest

```
1 #Making predictions on test dataset
2 rf_pred=rf.predict(X_test)
```

```
1 rf_pred.shape

(154,)
```

## Model Evaluation

### Train Score & Test Score

```
1 # Train score & Test score of logistic regression
2 from sklearn.metrics import accuracy_score
3 print("Train Accuracy of Logistic Regression",ln.score(X_train,y_train)*100)
```

```

3 print("Train Accuracy of Logistic Regression",lr.score(X_train,y_train)*100)
4 print("Accuracy (test) score of logistic Regression",lr.score(X_test, y_test)*100)
5 print("Accuracy (Test) score of Logistic Regression",accuracy_score(y_test, lr_pred)*100)

```

```

Train Accuracy of Logistic Regression 77.36156351791531
Accuracy (test) score of logistic Regression 77.27272727272727
Accuracy (Test) score of Logistic Regression 77.27272727272727

```

```

1 # Train score & Test score of KNN
2 print("Train Accuracy of KNN",knn.score(X_train,y_train)*100)
3 print("Accuracy (Test) score of KNN",knn.score(X_test, y_test)*100)
4 print("Accuracy score of KNN",accuracy_score(y_test, knn_pred)*100)

```

```

Train Accuracy of KNN 81.10749185667753
Accuracy (Test) score of KNN 74.67532467532467
Accuracy score of KNN 74.67532467532467

```

```

1 # Train score & Test score of Naive-Bayes
2 print("Train Accuracy of Naive Bayes",nb.score(X_train,y_train)*100)
3 print("Accuracy (Test) score of Naive Bayes",nb.score(X_test, y_test)*100)
4 print("Accuracy score of Naive Bayes",accuracy_score(y_test,nb_pred)*100)

```

```

Train Accuracy of Naive Bayes 74.2671009771987
Accuracy (Test) score of Naive Bayes 74.02597402597402
Accuracy score of Naive Bayes 74.02597402597402

```

```

1 # Train score & Test score of SVM
2 print("Train Accuracy of SVM",sv.score(X_train, y_train)*100)
3 print("Accuracy (Test) score of SVM",sv.score(X_test, y_test)*100)
4 print("Accuracy score of SVM",accuracy_score(y_test,sv_pred)*100)

```

```

Train Accuracy of SVM 81.92182410423453
Accuracy (Test) score of SVM 83.11688311688312
Accuracy score of SVM 83.11688311688312

```

```

1 # Train score & Test score of Decesion Tree
2 print("Train Accuracy of Decesion Tree",dt.score(X_train,y_train)*100)
3 print("Accuracy(Test) score of Decesion Tree",dt.score(X_test,y_test)*100)
4 print("Accuracy score of Decesion Tree",accuracy_score(y_test,dt_pred)*100)

```

```

Train Accuracy of Decesion Tree 100.0
Accuracy(Test) score of Decesion Tree 79.87012987012987
Accuracy score of Decesion Tree 79.87012987012987

```

```

1 # Train score & Test score of Random Forest
2 print("Train Accuracy of Random Forest",rf.score(X_train,y_train)*100)
3 print("Accuracy (Test) score of Random Forest",rf.score(X_test,y_test)*100)
4 print("Accuracy score of Random Forest",accuracy_score(y_test,rf_pred)*100)

```

```

Train Accuracy of Random Forest 100.0
Accuracy (Test) score of Random Forest 81.16883116883116
Accuracy score of Random Forest 81.16883116883116

```

## Confusion Matrix


### Confusion Matrix of "Logistic Regression"

```
1 from sklearn.metrics import classification_report, confusion_matrix
2 # confusion Matrix of Logistic Regression
3 cm=confusion_matrix(y_test,lr_pred)
4 cm
```

```
array([[86, 11],
       [24, 33]])
```

```
1 sns.heatmap(confusion_matrix(y_test,lr_pred),annot=True,fmt="d")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f7c1d14b4d0>
```



```
1 TN = cm[0,0]
2 FP = cm[0,1]
3 FN = cm[1,0]
4 TP = cm[1,1]
```

```
1 TN, FP, FN, TP
```

```
(86, 11, 24, 33)
```

```
1 #####
```

```
1 print('Classification Report of Logistic Regression: \n',classification_report(y_test,lr_pred))
```

```
Classification Report of Logistic Regression:
              precision    recall  f1-score   support

     0       0.7818      0.8866      0.8309         97
     1       0.7500      0.5789      0.6535         57

 accuracy          0.7727         154
 macro avg       0.7659      0.7328      0.7422         154
```



weighted avg	0.7700	0.7727	0.7652	154
--------------	--------	--------	--------	-----

```

1 # Making the Confusion Matrix Of Logistic Regression
2 from sklearn.metrics import classification_report, confusion_matrix
3 from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve
4 cm = confusion_matrix(y_test, lr_pred)
5
6 print('TN - True Negative {}'.format(cm[0,0]))
7 print('FP - False Positive {}'.format(cm[0,1]))
8 print('FN - False Negative {}'.format(cm[1,0]))
9 print('TP - True Positive {}'.format(cm[1,1]))
10 print('Accuracy Rate: {}'.format(np.divide(np.sum([cm[0,0],cm[1,1]]),np.sum(cm))*100))
11 print('Misclassification Rate: {}'.format(np.divide(np.sum([cm[0,1],cm[1,0]]),np.sum(cm)

```

```

    TN - True Negative 86
    FP - False Positive 11
    FN - False Negative 24
    TP - True Positive 33
    Accuracy Rate: 77.27272727272727
    Misclassification Rate: 22.727272727272727

```

```

1 77.27272727272727+22.727272727272727

```

```

    100.0

```

```

1 import matplotlib.pyplot as plt
2 plt.clf()
3 plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Wistia)
4 classNames = ['0','1']
5 plt.title('Confusion Matrix of Logistic Regression')
6 plt.ylabel('Actual(true) values')
7 plt.xlabel('Predicated values')
8 tick_marks = np.arange(len(classNames))
9 plt.xticks(tick_marks, classNames, rotation=45)
10 plt.yticks(tick_marks, classNames)
11 s = [['TN','FP'],['FN','TP']]
12 for i in range(2):
13     for j in range(2):
14         plt.text(j,i,str(s[i][j])+"="+str(cm[i][j]))
15 plt.show()

```

Confusion Matrix of Logistic Regression



```
1 pd.crosstab(y_test, lr_pred, margins=False)
```

col_0	0	1
Outcome		
0	86	11
1	24	33



```
1 pd.crosstab(y_test, lr_pred, margins=True)
```

col_0	0	1	All
Outcome			
0	86	11	97
1	24	33	57
All	110	44	154

```
1 pd.crosstab(y_test, lr_pred, rownames=['Actual values'], colnames=['Predicted values'],
```

Predicted values	0	1	All
Actual values			
0	86	11	97
1	24	33	57
All	110	44	154

```
1 TP,FP
```

(33, 11)

```
1 Precision=TP/(TP+FP)
```

```
2 Precision
```

0.75

```
1 33/(33+11)
```

0.75

```
1 # print precision score
```

```
2
```

```
3 precision_Score = TP/float(TP+FP)*100
```

```
4 print('Precision score : {0:0.4f}'.format(precision_Score))
```

```
4 print( Precision score : {0:0.4f} .format(precision_score))
```

```
Precision score : 75.0000
```

```
1 from sklearn.metrics import precision_score
2 print("precision Score is:", precision_score(y_test,lr_pred,average='micro')*100)
3 print("Mircro Average precision Score is:",precision_score(y_test,lr_pred,average='micr
4 print("Marcro Average precision Score is:",precision_score(y_test,lr_pred,average='macr
5 print("Weighted Average precision Score is:",precision_score(y_test,lr_pred,average='we
6 print("precision Score on Non weighted score is:", precision_score(y_test,lr_pred,avera
```

```
precision Score is: 77.27272727272727
Mircro Average precision Score is: 77.27272727272727
Marcro Average precision Score is: 76.5909090909091
Weighted Average precision Score is: 77.00413223140497
precision Score on Non weighted score is: [78.18181818 75.          ]
```

```
1 print('Classification Report of Logistic Regression: \n',classification_report(y_test,l
```

```
Classification Report of Logistic Regression:
              precision    recall  f1-score   support

         0       0.7818      0.8866      0.8309         97
         1       0.7500      0.5789      0.6535         57

 accuracy          0.7727         154
 macro avg          0.7659      0.7328      0.7422         154
weighted avg          0.7700      0.7727      0.7652         154
```

```
1 recall_score=TP/float(TP+FN)*100
2 print('recall_score',recall_score)
```

```
recall_score 57.89473684210527
```

```
1 TP,FN
```

```
(33, 24)
```

```
1 33/(33+24)
```

```
0.5789473684210527
```

```
1 from sklearn.metrics import recall_score
2 print('Recall or Sensitivity_score:',recall_score(y_test,lr_pred)*100)
```

```
Recall or Sensitivity_score: 57.89473684210527
```

```
1 print("Mircro Average Recall Score is:", recall_score(y_test,lr_pred,average='micro')*1
2 print("Marcro Average Recall Score is:", recall_score(y_test,lr_pred,average='macro')*1
3 print("Weighted Average Recall Score is:",recall_score(y_test,lr_pred,average='weighted
4 print("Recall Score on Non weighted score is:", recall_score(y_test,lr_pred,average=Non
```

Micro Average Recall Score is: 77.27272727272727  
 Macro Average Recall Score is: 73.27726532826912  
 Weighted Average Recall Score is: 77.27272727272727  
 Recall Score on Non weighted score is: [88.65979381 57.89473684]

```
1 print('Classification Report of logistic Regression: \n',classification_report(y_test,1
```

```
Classification Report of logistic Regression:
              precision    recall  f1-score   support

     0       0.7818      0.8866      0.8309         97
     1       0.7500      0.5789      0.6535         57

 accuracy          0.7727         154
 macro avg       0.7659      0.7328      0.7422         154
 weighted avg    0.7700      0.7727      0.7652         154
```

## False Positive Rate(FPR)

```
1 FPR=FP/float(FP+TN)*100
2 print('False Positive Rate : {0:0.4f}'.format(FPR))
```

False Positive Rate : 11.3402

```
1 FP,TN
```

(11, 86)

```
1 11/(11+86)
```

0.1134020618556701

## Specificity

```
1 specificity=TN/(TN+FP)*100
2 print('Specificity :{0:0.4f}'.format(specificity))
```

Specificity :88.6598

## F1-Score

```
1 from sklearn.metrics import f1_score
2 print('f1_score of macro :',f1_score(y_test, lr_pred)*100)
```

f1\_score of macro : 65.34653465346535

```
1 print("Micro Average F1_Score is:",f1_score(y_test,lr_pred,average='micro')*100)
2 print("macro Average F1_Score is:",f1_score(y_test,lr_pred,average='macro')*100)
```

```
3 print("Weighted Average F1_score is:",f1_score(y_test,lr_pred,average='weighted')*100)
4 print("F1_Score on Non weighted score is:",f1_score(y_test,lr_pred,average=None)*100)
```

```
Mircro Average F1_Score is: 77.27272727272727
marcro Average F1_Score is: 74.21916104653944
Weighted Average F1_score is: 76.52373933045479
F1_Score on Non weighted score is: [83.09178744 65.34653465]
```

### \*Classification Report of Logistic Regression \*

```
1 from sklearn.metrics import classification_report
2 print('Classification Report of logistic Regression: \n', classification_report(y_test,
```

```
Classification Report of logistic Regression:
              precision    recall  f1-score   support

         0           0.7818        0.8866        0.8309          97
         1           0.7500        0.5789        0.6535          57


 accuracy                   0.7727          154
 macro avg           0.7659        0.7328        0.7422          154
 weighted avg        0.7700        0.7727        0.7652          154
```

### ROC Curve & ROC AUC

```
1 # Area Under Curve
2 auc = roc_auc_score(y_test,lr_pred)
3 print("ROC AUC SCORE of logistic Regression is",auc)
```

```
ROC AUC SCORE of logistic Regression is 0.7327726532826913
```

```
1 fpr, tpr, thresholds = roc_curve(y_test, lr_pred)
2 plt.plot(fpr, tpr, color='orange', label='ROC')
3 plt.plot([0,1],[0,1], color='darkblue', linestyle='--',label='ROC curve (area = %0.2f)')
4 plt.xlabel('False Positive Rate')
5 plt.ylabel('True Positive Rate')
6 plt.title('Receiver Operating Characteristic(ROC) Curve of Logistic Regression')
7 plt.legend()
8 plt.grid()
9 plt.show()
```




## Confusion Matrix of "KNN"



```
1 sns.heatmap(confusion_matrix(y_test,knn_pred),annot=True,fmt="d")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f7c1d2fcd90>
```



```
1 TN = cm[0,0]
```

```
2 FP = cm[0,1]
```

```
3 FN = cm[1,0]
```

```
4 TP = cm[1,1]
```

```
1 TN, FP, FN, TP
```

```
(86, 11, 24, 33)
```

```
1 #####
```

```
1 # classification Report of KNN
```

```
2 print('Classification Report of KNN: \n', classification_report(y_test,knn_pred,digits=
```

```
Classification Report of KNN:
```

	precision	recall	f1-score	support
0	0.7736	0.8454	0.8079	97
1	0.6875	0.5789	0.6286	57
accuracy			0.7468	154
macro avg	0.7305	0.7122	0.7182	154
weighted avg	0.7417	0.7468	0.7415	154

```
1 # Making the confusion Matrix of KNN
```

```
2 from sklearn.metrics import classification_report, confusion_matrix
```

```
3 from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve
```

```

4 cm = confusion_matrix(y_test,knn_pred)
5
6
7 print('TN - True Negative: {}'.format(cm[0,0]))
8 print('FP - False Positive: {}'.format (cm[0,1]))
9 print('FN - False Negative: {}'.format(cm[1,0]))
10 print('TP - True Positive: {}'.format(cm[1,1]))
11 print('Accuracy Rate: {}'.format(np.divide(np.sum([cm[0,0],cm[0,1],cm[1,0],cm[1,1]]),np
12 print('Misclassification Rate: {}'.format(np.divide(np.sum([cm[0,1],cm[1,0]]),np.sum(cm

```

```

TN - True Negative: 82
FP - False Positive: 15
FN - False Negative: 24
TP - True Positive: 33
Accuracy Rate: 100.0
Misclassification Rate: 25.324675324675322

```

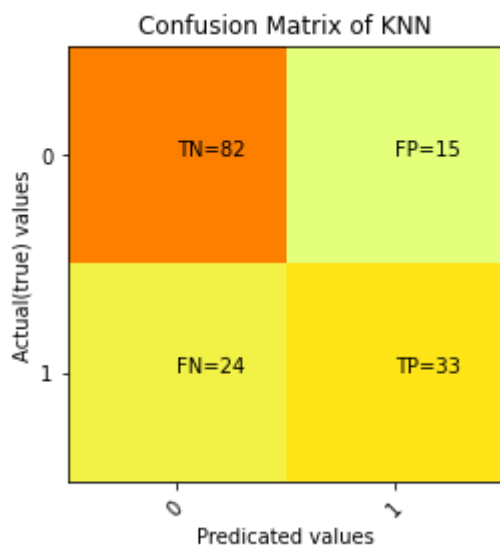
```
1 74.67532467532467+25.324675324675322
```

```
100.0
```

```

1 import matplotlib.pyplot as plt
2 plt.clf()
3 plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Wistia)
4 classNames = ['0','1']
5 plt.title('Confusion Matrix of KNN')
6 plt.ylabel('Actual(true) values')
7 plt.xlabel('Predicated values')
8 tick_marks = np.arange(len(classNames))
9 plt.xticks(tick_marks, classNames, rotation=45)
10 plt.yticks(tick_marks, classNames)
11 s = [['TN','FP'],['FN','TP']]
12 for i in range(2):
13     for j in range(2):
14         plt.text(j,i,str(s[i][j])+"="+str(cm[i][j]))
15 plt.show()

```



```
1 nd crossstab(y_test, knn_pred, margins=False)
```

```
1 pd.crosstab(y_test, knn_pred, margins=True)
```

	col_0	0	1
Outcome			
0	82	15	
1	24	33	

```
1 pd.crosstab(y_test, knn_pred, margins=True)
```

	col_0	0	1	All
Outcome				
0	82	15	97	
1	24	33	57	
All	106	48	154	

```
1 pd.crosstab(y_test, knn_pred, rownames=['Actual values'], colnames=['Predicted values'])
```

	Predicted values	0	1	All
Actual values				
0	82	15	97	
1	24	33	57	
All	106	48	154	

```
1 TP,FP
```

```
(33, 11)
```

```
1 Precision=TP/(TP+FP)
```

```
2 Precision
```

```
0.75
```

```
1 33/(33+15)
```

```
0.6875
```

```
1 # print precision score
```

```
2
```

```
3 precision_Score = TP/float(TP+FP)*100
```

```
4 print('Precision score : {0:0.4f}'.format(precision_Score))
```

```
Precision score : 75.0000
```



```

1 from sklearn.metrics import precision_score
2 print("precision Score is:", precision_score(y_test,knn_pred,average='micro')*100)
3 print("Mircro Average precision Score is:",precision_score(y_test,knn_pred,average='mic
4 print("Marco Average precision Score is:",precision_score(y_test,knn_pred,average='mac
5 print("Weighted Average precision Score is:",precision_score(y_test,knn_pred,average='w
6 print("precision Score on Non weighted score is:", precision_score(y_test,knn_pred,aver

```

```

precision Score is: 74.67532467532467
Mircro Average precision Score is: 74.67532467532467
Marco Average precision Score is: 73.05424528301887
Weighted Average precision Score is: 74.17223107081597
precision Score on Non weighted score is: [77.35849057 68.75      ]

```

```

1 print('Classification Report of KNN: \n',classification_report(y_test,knn_pred,digits=4

```

```

Classification Report of KNN:

```

	precision	recall	f1-score	support
0	0.7736	0.8454	0.8079	97
1	0.6875	0.5789	0.6286	57
accuracy			0.7468	154
macro avg	0.7305	0.7122	0.7182	154
weighted avg	0.7417	0.7468	0.7415	154

```

1 recall_score=TP/float(TP+FN)*100
2 print('recall_score',recall_score)

```

```

recall_score 57.89473684210527

```

```

1 TP,FN

```

```

(33, 24)

```

```

1 33/(33+24)

```

```

0.5789473684210527

```

```

1 from sklearn.metrics import recall_score
2 print('Recall or Sensitivity_score:',recall_score(y_test,knn_pred)*100)

```

```

Recall or Sensitivity_score: 57.89473684210527

```

```

1 print("Mircro Average Recall Score is:", recall_score(y_test,knn_pred,average='micro')*
2 print("Marco Average Recall Score is:", recall_score(y_test,knn_pred,average='macro')*
3 print("Weighted Average Recall Score is:",recall_score(y_test,knn_pred,average='weighte
4 print("Recall Score on Non weighted score is:", recall_score(y_test,knn_pred,average=No

```

```

Mircro Average Recall Score is: 74.67532467532467
Marco Average Recall Score is: 71.21540965816604
Weighted Average Recall Score is: 74.67532467532467
Recall Score on Non weighted score is: [84.53608247 57.89473684]


```

## ROC Curve & ROC AUC

```
1 # Area Under the Curve
2 auc = roc_auc_score(y_test, knn_pred)
3 print("ROC AUC SCORE of KNN is",auc)
```

ROC AUC SCORE of KNN is 0.7121540965816603

```
1 fpr, tpr, thresholds = roc_curve(y_test, knn_pred)
2 plt.plot(fpr, tpr, color='orange', label='ROC')
3 plt.plot([0,1],[0,1], color='darkblue', linestyle='--',label='ROC curve (area = %0.2f)')
4 plt.xlabel('False Positive Rate')
5 plt.ylabel('True Positive Rate')
6 plt.title('Receiver Operating Characteristic(ROC) Curve of KNN')
7 plt.legend()
8 plt.grid()
9 plt.show()
```



## Confusion Matrix of "Naive-Byes"

```
1 sns.heatmap(confusion_matrix(y_test, nb_pred),annot=True,fmt="d")
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f7c1d897c50>

```
1 TN = cm[0,0]
2 FP = cm[0,1]
3 FN = cm[1,0]
4 TP = cm[1,1]
```

```
1 TN, FP, FN,TP
```

(82, 15, 24, 33)

```
1 #####
```

```
1 # classification Report of Naive Byes
```

```
2 print('Classification Report of Naive Byes: \n', classification_report(y_test,nb_pred,d
```

```
Classification Report of Naive Byes:
              precision    recall  f1-score   support

     0       0.7879      0.8041      0.7959         97
     1       0.6545      0.6316      0.6429         57

 accuracy          0.7403         154
 macro avg       0.7212      0.7179      0.7194         154
weighted avg       0.7385      0.7403      0.7393         154
```

```
1 # Making the confusion Matrix of Naive Bayes
```

```
2 from sklearn.metrics import classification_report, confusion_matrix
```

```
3 from sklearn.metrics import accuracy_score, roc_auc_score,roc_curve
```

```
4 cm = confusion_matrix(y_test,nb_pred)
```

```
5
```

```
6
```

```
7 print('TN - True Negative: {}'.format(cm[0,0]))
```

```
8 print('FP - False Positive: {}'.format (cm[0,1]))
```

```
9 print('FN - False Negative: {}'.format(cm[1,0]))
```

```
10 print('TP - True Positive: {}'.format(cm[1,1]))
```

```
11 print('Accuracy Rate: {}'.format(np.divide(np.sum([cm[0,0],cm[0,1],cm[1,0],cm[1,1]]),np
```

```
12 print('Misclassification Rate: {}'.format(np.divide(np.sum([cm[0,1],cm[1,0]]),np.sum(cm
```

```
TN - True Negative: 78
FP - False Positive: 19
FN - False Negative: 21
TP - True Positive: 36
Accuracy Rate: 100.0
Misclassification Rate: 25.97402597402597
```

```
1 74.02597402597402+25.97402597402597
```


```
100.0
```

```
1 import matplotlib.pyplot as plt
```

```

2 plt.clt()
3 plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Wistia)
4 classNames = ['0','1']
5 plt.title('Confusion Matrix of Naive Byes')
6 plt.ylabel('Actual(true) values')
7 plt.xlabel('Predicated values')
8 tick_marks = np.arange(len(classNames))
9 plt.xticks(tick_marks, classNames, rotation=45)
10 plt.yticks(tick_marks, classNames)
11 s = [['TN','FP'], ['FN','TP']]
12 for i in range(2):
13     for j in range(2):
14         plt.text(j,i,str(s[i][j])+"="+str(cm[i][j]))
15 plt.show()

```



```
1 pd.crosstab(y_test, nb_pred, margins=False)
```

col_0	0	1
Outcome		
0	78	19
1	21	36

```
1 pd.crosstab(y_test, nb_pred, margins=True)
```

col_0	0	1	All
Outcome			
0	78	19	97
1	21	36	57
All	99	55	154

```
1 pd.crosstab(y_test, nb_pred, rownames=['Actual values'], colnames=['Predicted values'],
```

Predicted values	0	1	All
Actual values			
0	78	19	97
1	21	36	57
All	99	55	154

1 TP,FP

(44, 18)

1 Precision=TP/(TP+FP)

2 Precision

0.6875

1 36/(36+19)

0.6545454545454545

1 # print precision score

2

3 precision\_Score = TP/float(TP+FP)\*100

4 print('Precision score : {0:0.4f}'.format(precision\_Score))

Precision score : 68.7500

1 from sklearn.metrics import precision\_score

2 print("precision Score is:", precision\_score(y\_test,nb\_pred,average='micro')\*100)

3 print("Mircro Average precision Score is:",precision\_score(y\_test,nb\_pred,average='micr

4 print("Marcro Average precision Score is:",precision\_score(y\_test,nb\_pred,average='macr

5 print("Weighted Average precision Score is:",precision\_score(y\_test,nb\_pred,average='we

6 print("precision Score on Non weighted score is:", precision\_score(y\_test,nb\_pred,avera

precision Score is: 74.02597402597402

Mircro Average precision Score is: 74.02597402597402

Marcro Average precision Score is: 72.12121212121212

Weighted Average precision Score is: 73.85281385281385

precision Score on Non weighted score is: [78.78787879 65.45454545]

1 # classification Report of Naive Bayes

2 print('Classification Report of Naive Bayes: \n', classification\_report(y\_test,nb\_pred,

Classification Report of Naive Bayes:

precision recall f1-score support

0 0.7879 0.8041 0.7959 97

1 0.6545 0.6316 0.6429 57

accuracy 0.7403 154

macro avg	0.7212	0.7179	0.7194	154
weighted avg	0.7385	0.7403	0.7393	154

```
1 recall_score=TP/float(TP+FN)*100
2 print('recall_score',recall_score)
```

```
recall_score 57.89473684210527
```

```
1 TP,FN
```

```
(33, 24)
```

```
1 36/(36+21)
```

```
0.631578947368421
```

```
1 from sklearn.metrics import recall_score
2 print('Recall or Sensitivity_score:',recall_score(y_test,nb_pred)*100)
```


```
Recall or Sensitivity_score: 63.1578947368421
```

## ROC Curve & ROC AUC

```
1 # Area Under the Curve
2 auc = roc_auc_score(y_test, nb_pred)
3 print("ROC AUC SCORE of Naive Bayes is",auc)
```

```
ROC AUC SCORE of Naive Bayes is 0.7178513293543136
```


```
1 fpr, tpr, thresholds = roc_curve(y_test, nb_pred)
2 plt.plot(fpr, tpr, color='orange', label='ROC')
3 plt.plot([0,1],[0,1], color='darkblue', linestyle='--',label='ROC curve (area = %0.2f)')
4 plt.xlabel('False Positive Rate')
5 plt.ylabel('True Positive Rate')
6 plt.title('Receiver Operating Characteristic(ROC) Curve of Naive Bayes')
7 plt.legend()
8 plt.grid()
9 plt.show()
```



## Confusion Matrix of "SVM"

1 `sns.heatmap(confusion_matrix(y_test, sv_pred),annot=True,fmt="d")`

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f7c1d9dc210>



```
1 TN = cm[0,0]
2 FP = cm[0,1]
3 FN = cm[1,0]
4 TP = cm[1,1]
```

```
1 TN, FP, FN,TP
```

```
(78, 19, 21, 36)
```

```
1 #####
```

```
1 # classification Report of SVM
```

```
2 print('Classification Report of SVM: \n', classification_report(y_test,sv_pred,digits=4
```

Classification Report of SVM:

	precision	recall	f1-score	support
0	0.8198	0.9381	0.8750	97
1	0.8605	0.6491	0.7400	57
accuracy			0.8312	154
macro avg	0.8401	0.7936	0.8075	154
weighted avg	0.8349	0.8312	0.8250	154

```
1 # Making the confusion Matrix of SVM
```

```
2 from sklearn.metrics import classification_report, confusion_matrix
```

```
3 from sklearn.metrics import accuracy_score, roc_auc_score,roc_curve
```

```

4 cm = confusion_matrix(y_test,sv_pred)
5
6
7 print('TN - True Negative: {}'.format(cm[0,0]))
8 print('FP - False Positive: {}'.format (cm[0,1]))
9 print('FN - False Negative: {}'.format(cm[1,0]))
10 print('TP - True Positive: {}'.format(cm[1,1]))
11 print('Accuracy Rate: {}'.format(np.divide(np.sum([cm[0,0],cm[0,1],cm[1,0],cm[1,1]]),np
12 print('Misclassification Rate: {}'.format(np.divide(np.sum([cm[0,1],cm[1,0]]),np.sum(cm

```

```

TN - True Negative: 91
FP - False Positive: 6
FN - False Negative: 20
TP - True Positive: 37
Accuracy Rate: 100.0
Misclassification Rate: 16.883116883116884

```

```

1 83.11688311688312+16.883116883116884

```

```


100.0

```

```

1 import matplotlib.pyplot as plt
2 plt.clf()
3 plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Wistia)
4 classNames = ['0','1']
5 plt.title('Confusion Matrix of Naive Byes')
6 plt.ylabel('Actual(true) values')
7 plt.xlabel('Predicated values')
8 tick_marks = np.arange(len(classNames))
9 plt.xticks(tick_marks, classNames, rotation=45)
10 plt.yticks(tick_marks, classNames)
11 s = [['TN','FP'],['FN','TP']]
12 for i in range(2):
13     for j in range(2):
14         plt.text(j,i,str(s[i][j])+"="+str(cm[i][j]))
15 plt.show()

```



```

1 pd.crosstab(y_test, sv_pred, margins=False)

```



col_0	0	1
Outcome		
0	91	6
1	20	37

```
1 pd.crosstab(y_test, sv_pred, margins=True)
```

col_0	0	1	All
Outcome			
0	91	6	97
1	20	37	57
All	111	43	154

```
1 pd.crosstab(y_test, sv_pred, rownames=['Actual values'], colnames=['Predicted values'],
```

Predicted values	0	1	All
Actual values			
0	91	6	97
1	20	37	57
All	111	43	154

```
1 TP, FP
```

```
(36, 19)
```

```
1 Precision=TP/(TP+FP)
```

```
2 Precision
```

```
0.6545454545454545
```

```
1 37/(37+6)
```

```
0.8604651162790697
```

```
1 #print 'Precision score'
```

```
2
```

```
3 precision_Score = TP/float(TP+FP)*100
```

```
4 print('Precision score : {0:0.4f}'.format(precision_Score))
```

```
Precision score : 65.4545
```

```

1 from sklearn.metrics import precision_score
2 print("precision Score is:", precision_score(y_test,sv_pred,average='micro')*100)
3 print("Mircro Average precision Score is:",precision_score(y_test,sv_pred,average='micr
4 print("Marcro Average precision Score is:",precision_score(y_test,sv_pred,average='macr
5 print("Weighted Average precision Score is:",precision_score(y_test,sv_pred,average='we
6 print("precision Score on Non weighted score is:", precision_score(y_test,sv_pred,avera

```

```

precision Score is: 83.11688311688312
Mircro Average precision Score is: 83.11688311688312
Marcro Average precision Score is: 84.01424680494446
Weighted Average precision Score is: 83.48638581196721
precision Score on Non weighted score is: [81.98198198 86.04651163]

```

```

1 # classification Report of SVM
2 print('Classification Report of SVM: \n', classification_report(y_test,sv_pred,digits=4

```

```

Classification Report of SVM:

```

	precision	recall	f1-score	support
0	0.8198	0.9381	0.8750	97
1	0.8605	0.6491	0.7400	57
accuracy			0.8312	154
macro avg	0.8401	0.7936	0.8075	154
weighted avg	0.8349	0.8312	0.8250	154

```

1 recall_score=TP/float(TP+FN)*100
2 print('recall_score',recall_score)

```

```

recall_score 63.1578947368421

```

```

1 TP, FN

```

```

(36, 21)

```

```

1 from sklearn.metrics import recall_score
2 print('Recall or Sensitivity_score:',recall_score(y_test,sv_pred)*100)

```

```

Recall or Sensitivity_score: 64.91228070175438

```

```

1 37/(37+20)

```

```

0.6491228070175439

```

## ROC Curve & ROC AUC

```

1 # Area Under the Curve
2 auc = roc_auc_score(y_test, sv_pred)
3 print("ROC AUC SCORE of SVM is",auc)

```

```


ROC AUC SCORE of SVM is 0.7936335684572255

```

```

1 fpr, tpr, thresholds = roc_curve(y_test, sv_pred)
2 plt.plot(fpr, tpr, color='orange', label='ROC')
3 plt.plot([0,1],[0,1], color='darkblue', linestyle='--',label='ROC curve (area = %0.2f)')
4 plt.xlabel('False Positive Rate')
5 plt.ylabel('True Positive Rate')
6 plt.title('Receiver Operating Characteristic(ROC) Curve of SVM')
7 plt.legend()
8 plt.grid()
9 plt.show()


```



## Confusion Matrix of "Decision Tree"

```
1 sns.heatmap(confusion_matrix(y_test, dt_pred),annot=True,fmt="d")
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f7c1d407e90>



```

1 TN = cm[0,0]
2 FP = cm[0,1]
3 FN = cm[1,0]
4 TP = cm[1,1]

```

```
1 TN, FP, FN, TP
```

(91, 6, 20, 37)

```
1 #####
```

```
1 # classification Report of Decesion Tree
```

```
2 print('Classification Report of Decesion Tree: \n', classification_report(y_test,dt_pre
```

```
Classification Report of Decesion Tree:
```

	precision	recall	f1-score	support
0	0.8587	0.8144	0.8360	97
1	0.7097	0.7719	0.7395	57
accuracy			0.7987	154
macro avg	0.7842	0.7932	0.7877	154
weighted avg	0.8035	0.7987	0.8003	154

```
1 # Making the confusion Matrix of Decesion Tree
```

```
2 from sklearn.metrics import classification_report, confusion_matrix
```

```
3 from sklearn.metrics import accuracy_score, roc_auc_score,roc_curve
```

```
4 cm = confusion_matrix(y_test,dt_pred)
```

```
5
```

```
6
```

```
7 print('TN - True Negative: {}'.format(cm[0,0]))
```

```
8 print('FP - False Positive:{}'.format (cm[0,1]))
```

```
9 print('FN - False Negative:{}'.format(cm[1,0]))
```

```
10 print('TP - True Positive:{}'.format(cm[1,1]))
```

```
11 print('Accuracy Rate: {}'.format(np.divide(np.sum([cm[0,0],cm[0,1],cm[1,0],cm[1,1]]),np
```

```
12 print('Misclassification Rate: {}'.format(np.divide(np.sum([cm[0,1],cm[1,0]]),np.sum(cm
```

```
TN - True Negative: 79
```

```
FP - False Positive:18
```

```
FN - False Negative:13
```

```
TP - True Positive:44
```

```
Accuracy Rate: 100.0
```

```
Misclassification Rate: 20.12987012987013
```

```
1 78.57142857142857+21.428571428571427
```

```
100.0
```

```
1 import matplotlib.pyplot as plt
```

```
2 plt.clf()
```

```
3 plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Wistia)
```

```
4 classNames = ['0','1']
```

```
5 plt.title('Confusion Matrix of Decision Tree')
```

```
6 plt.ylabel('Actual(true) values')
```

```
7 plt.xlabel('Predicated values')
```

```
8 tick_marks = np.arange(len(classNames))
```

```
9 plt.xticks(tick_marks, classNames, rotation=45)
```


```
10 plt.yticks(tick_marks, classNames)
```

```
11 s = [['TN','FP'],['FN','TP']]
```

```

12 for i in range(2):
13     for j in range(2):
14         plt.text(j,i,str(s[i][j])+"="+str(cm[i][j]))
15 plt.show()

```



```
1 pd.crosstab(y_test, dt_pred, margins=False)
```

col_0	0	1
Outcome		
0	79	18
1	13	44

```
1 pd.crosstab(y_test, dt_pred, margins=True)
```

col_0	0	1	All
Outcome			
0	79	18	97
1	13	44	57
All	92	62	154

```
1 pd.crosstab(y_test, dt_pred, rownames=['Actual values'], colnames=['Predicted values'],
```

Predicted values	0	1	All
Actual values			
0	79	18	97
1	13	44	57
All	92	62	154

```
1 TP, FP
```

```
(37, 6)
```

```
1 Precision=TP/(TP+FP)
```

```
2 Precision
```

```
0.8604651162790697
```

```
1 42/(42+17)
```

```
0.711864406779661
```

```
1 #print precision score
```

```
2
```

```
3 precision_Score = TP/float(TP+FP)*100
```

```
4 print('Precision score : {0:0.4f}'.format(precision_Score))
```

```
Precision score : 86.0465
```

```
1 from sklearn.metrics import precision_score
```

```
2 print("precision Score is:", precision_score(y_test,dt_pred,average='micro')*100)
```

```
3 print("Micro Average precision Score is:",precision_score(y_test,dt_pred,average='micro'))
```

```
4 print("Macro Average precision Score is:",precision_score(y_test,dt_pred,average='macro'))
```

```
5 print("Weighted Average precision Score is:",precision_score(y_test,dt_pred,average='weighted'))
```

```
6 print("precision Score on Non weighted score is:", precision_score(y_test,dt_pred,average='weighted'))
```

```
precision Score is: 79.87012987012987
```

```
Micro Average precision Score is: 79.87012987012987
```

```
Macro Average precision Score is: 78.41865357643759
```

```
Weighted Average precision Score is: 80.35395530136064
```

```
precision Score on Non weighted score is: [85.86956522 70.96774194]
```

```
1 # classification Report of Decesion Tree
```

```
2 print('Classification Report of Decesion Tree: \n', classification_report(y_test,dt_pre
```

```
Classification Report of Decesion Tree:
```

	precision	recall	f1-score	support
0	0.8587	0.8144	0.8360	97
1	0.7097	0.7719	0.7395	57
accuracy			0.7987	154
macro avg	0.7842	0.7932	0.7877	154
weighted avg	0.8035	0.7987	0.8003	154

```
1 recall_score=TP/float(TP+FN)*100
```

```
2 print('recall_score',recall_score)
```

```
recall_score 64.91228070175438
```

```
1 TP, FN
```

```
1 TP, FN
```

```
(37, 20)
```

```
1 from sklearn.metrics import recall_score
2 print('Recall or Sensitivity_score:', recall_score(y_test, dt_pred)*100)
```

```
Recall or Sensitivity_score: 77.19298245614034
```

```
1 42/(42+15)
```


```
0.7368421052631579
```

## ROC Curve & ROC AUC

```
1 # Area Under the Curve
2 auc = roc_auc_score(y_test, dt_pred)
3 print("ROC AUC SCORE of Decesion Tree is", auc)
```

```
ROC AUC SCORE of Decesion Tree is 0.7931814071260626
```


```
1 fpr, tpr, thresholds = roc_curve(y_test, dt_pred)
2 plt.plot(fpr, tpr, color='orange', label='ROC')
3 plt.plot([0,1],[0,1], color='darkblue', linestyle='--', label='ROC curve (area = %0.2f)')
4 plt.xlabel('False Positive Rate')
5 plt.ylabel('True Positive Rate')
6 plt.title('Receiver Operating Characteristic(ROC) Curve of Decesion Tree')
7 plt.legend()
8 plt.grid()
9 plt.show()
```



## Confusion Matrix of "Random Forest"

```
1 sns.heatmap(confusion_matrix(y_test, rf_pred), annot=True, fmt="d")
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f7c1d7b8c90>



```
1 TN = cm[0,0]
2 FP = cm[0,1]
3 FN = cm[1,0]
4 TP = cm[1,1]
```

```
1 TN, FP, FN,TP
```

```
(87, 10, 19, 38)
```

```
1 ####
```

```
1 # classification Report of Random Forest
2 print('Classification Report of Random Forest: \n', classification_report(y_test,rf_pre
```

```
Classification Report of Random Forest:
              precision    recall  f1-score   support

     0       0.8208      0.8969      0.8571        97
     1       0.7917      0.6667      0.7238        57

 accuracy          0.8117        154
 macro avg       0.8062      0.7818      0.7905        154
 weighted avg    0.8100      0.8117      0.8078        154
```

```
1 # Making the confusion Matrix of Random Forest
2 from sklearn.metrics import classification_report, confusion_matrix
3 from sklearn.metrics import accuracy_score, roc_auc_score,roc_curve
4 cm = confusion_matrix(y_test,rf_pred)
5
6
7 print('TN - True Negative: {}'.format(cm[0,0]))
8 print('FP - False Positive: {}'.format (cm[0,1]))
9 print('FN - False Negative: {}'.format(cm[1,0]))
10 print('TP - True Positive: {}'.format(cm[1,1]))
11 print('Accuracy Rate: {}'.format(np.divide(np.sum([cm[0,0],cm[0,1],cm[1,0],cm[1,1]]),np
12 print('Misclassification Rate: {}'.format(np.divide(np.sum([cm[0,1],cm[1,0]]),np.sum(cm
```



TN - True Negative: 87  
 FP - False Positive: 10  
 FN - False Negative: 19  
 TP - True Positive: 38  
 Accuracy Rate: 100.0  
 Misclassification Rate: 18.83116883116883

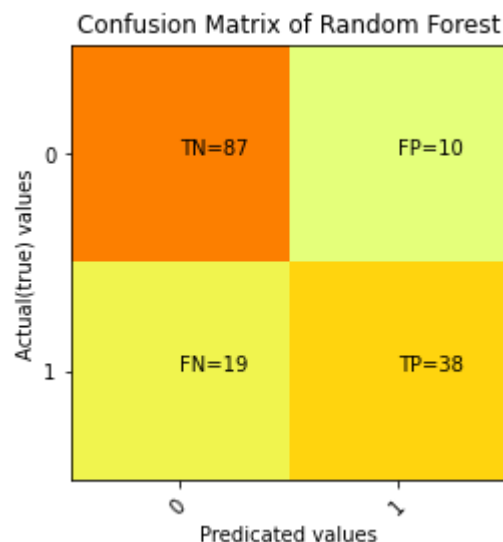
```
1 79.22077922077922+20.77922077922078
```

100.0

```

1 import matplotlib.pyplot as plt
2 plt.clf()
3 plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Wistia)
4 classNames = ['0', '1']
5 plt.title('Confusion Matrix of Random Forest')
6 plt.ylabel('Actual(true) values')
7 plt.xlabel('Predicated values')
8 tick_marks = np.arange(len(classNames))
9 plt.xticks(tick_marks, classNames, rotation=45)
10 plt.yticks(tick_marks, classNames)
11 s = [['TN', 'FP'], ['FN', 'TP']]
12 for i in range(2):
13     for j in range(2):
14         plt.text(j,i,str(s[i][j])+"="+str(cm[i][j]))
15 plt.show()

```



```
1 pd.crosstab(y_test, rf_pred, margins=False)
```

col_0	0	1
Outcome		
0	87	10
1	19	38

```
1 pd.crosstab(y_test, rf_pred, margins=True)
```

col_0	0	1	All
<b>Outcome</b>			
<b>0</b>	87	10	97
<b>1</b>	19	38	57
<b>All</b>	106	48	154

```
1 pd.crosstab(y_test, rf_pred, rownames=['Actual values'], colnames=['Predicted values'],
```

Predicted values	0	1	All
<b>Actual values</b>			
<b>0</b>	87	10	97
<b>1</b>	19	38	57
<b>All</b>	106	48	154

```
1 TP, FP
```

```
(38, 10)
```

```
1 Precision=TP/(TP+FP)
```

```
2 Precision
```

```
0.7916666666666666
```

```
1 38/(38+10)
```

```
0.7916666666666666
```

```
1 #print precision score
```

```
2
```

```
3 precision_Score = TP/float(TP+FP)*100
```

```
4 print('Precision score : {0:0.4f}'.format(precision_Score))
```

```
Precision score : 79.1667
```

```
1 from sklearn.metrics import precision_score
```

```
2 print("precision Score is:", precision_score(y_test,dt_pred,average='micro')*100)
```

```
3 print("Mircro Average precision Score is:",precision_score(y_test,dt_pred,average='micr
```

```
4 print("Marcro Average precision Score is:",precision_score(y_test,dt_pred,average='macr
```

```
5 print("Weighted Average precision Score is:",precision_score(y_test,dt_pred,average='we
```

```
6 print("precision Score on Non weighted score is:", precision_score(y_test,dt_pred,avera
```

```
precision Score is: 79.87012987012987
```

```
Mircro Average precision Score is: 79.87012987012987
```

```
Marcro Average precision Score is: 78.41865357643759
```

Weighted Average precision Score is: 80.35395530136064  
 precision Score on Non weighted score is: [85.86956522 70.96774194]

```
1 # classification Report of Random Forest
2 print('Classification Report of Random Forest: \n', classification_report(y_test,rf_pre
```

```
Classification Report of Random Forest:
              precision    recall  f1-score   support

     0       0.8208      0.8969      0.8571        97
     1       0.7917      0.6667      0.7238        57

 accuracy          0.8117        154
 macro avg       0.8062      0.7818      0.7905        154
 weighted avg    0.8100      0.8117      0.8078        154
```

```
1 recall_score=TP/float(TP+FN)*100
2 print('recall_score',recall_score)
```

```
recall_score 66.66666666666666
```

```
1 TP, FN
```

```
(38, 19)
```

```
1 from sklearn.metrics import recall_score
2 print('Recall or Sensitivity_score:',recall_score(y_test,rf_pred)*100)
```

```
Recall or Sensitivity_score: 66.66666666666666
```

```
1 39/(39+19)
```


```
0.6724137931034483
```

## ROC Curve & ROC AUC

```
1 # Area Under the Curve
2 auc = roc_auc_score(y_test, rf_pred)
3 print("ROC AUC SCORE of Random Forest is",auc)
```

```
ROC AUC SCORE of Random Forest is 0.781786941580756
```

```
1 fpr, tpr, thresholds = roc_curve(y_test, rf_pred)
2 plt.plot(fpr, tpr, color='orange', label='ROC')
3 plt.plot([0,1],[0,1], color='darkblue', linestyle='--',label='ROC curve (area = %0.2f)')
4 plt.xlabel('False Positive Rate')
5 plt.ylabel('True Positive Rate')
6 plt.title('Receiver Operating Characteristic(ROC) Curve of Random Forest')
7 plt.legend()
8 plt.grid()
9 plt.show()
```



```

1 # Confusion Matrix function
2
3 def conf_mtx(y_act,y_pred):
4     cm=metrics.confusion_matrix(y_act, y_pred, labels=[1, 0])
5     df_cm = pd.DataFrame(cm, index = [i for i in ["Diabetic","Non-Diabetic"]],
6                           columns = [i for i in ["Predict Diabetic","Predict Non-Diabetic"]])
7     plt.figure(figsize = (6,6))
8     plt.title("Confusion Matrix")
9     sns.heatmap(df_cm, annot=True ,fmt='g')
10
11     Score_Accuracy = "%.2f%%" %(metrics.accuracy_score(y_act,y_pred)*100)
12     Score_Recall = "%.2f%%" %(metrics.recall_score(y_act,y_pred)*100)
13     Score_Precision = "%.2f%%" %(metrics.precision_score(y_act,y_pred)*100)
14
15     print("Model Accuracy Score: " + Score_Accuracy)
16     print("Model Recall Score: " + Score_Recall)
17     print("Model Precision Score: " + Score_Precision)
18
19     return Score_Accuracy,Score_Recall,Score_Precision

```

```

1 # Prepare an empty summary dataframe to append the data of the various models for compa
2 summary = pd.DataFrame(columns=('Model', 'Training Accuracy', 'Test Accuracy Score','Te
3                               'Test Precision Score', 'AUC'))

```

```

1 # For building a function for performing ML algos testing
2
3 def ML_test(Mdl,Param_grid):
4     if bool(Param_grid):
5         Mdl = GridSearchCV(Mdl,Param_grid,cv=10)
6         Mdl.fit(X_train_sm,y_train_sm)
7         Mdl_params = Mdl.best_params_
8         Mdl_train_sc = Mdl.cv_results_['mean_test_score'].mean()
9         Mdl_test_sc = Mdl.score(X_test_scaled,y_test)
10        probas = Mdl.predict_proba(X_test_scaled)
11
12        print("Best fit parameter is: " + str(Mdl_params))

```

```

13
14     else:
15         Mdl = Mdl
16         Mdl.fit(X_train_sm,y_train_sm)
17         Mdl_train_sc = round(Mdl.score(X_train_sm,y_train_sm),4)
18         Mdl_test_sc = round(Mdl.score(X_test_scaled,y_test),4)
19         probas = Mdl.predict_proba(X_test_scaled)
20
21     y_pred = Mdl.predict(X_test_scaled)
22
23     print("Training score is: " + str(Mdl_train_sc))
24     print("Test Mean score is: " + str(Mdl_test_sc))
25
26     Score_Accuracy,Score_Recall,Score_Precision = conf_mtx(y_test,y_pred)
27     Mdl_train_sc = "%.2f%%" % (Mdl_train_sc*100)
28
29     # Calculating AUC
30     fpr, tpr, thresholds = roc_curve(y_test, probas[:, 1])
31     roc_auc = round(auc(fpr, tpr),4)
32     print("Area under the ROC curve : " + str(roc_auc))
33
34     return Mdl_train_sc, Score_Accuracy, Score_Recall, Score_Precision, roc_auc

```

```

1 # Scaling the x training and testing dataset
2 scaler = preprocessing.StandardScaler().fit(X_train)
3
4 X_train_scaled = scaler.transform(X_train)
5 X_test_scaled = scaler.transform(X_test)

```

```

1 # Logistic Regression Model
2
3 Mdl_LogReg = LogisticRegression(solver="liblinear")
4
5 model_name = "LogisticRegression"
6 Mdl_train_sc = "77.36156351791531"
7 Score_Accuracy = "77.27272727272727"
8 Score_Recall = "57.89473684210527"
9 Score_Precision = "75.27272727272727"
10 roc_auc = "0.7327726532826913"
11
12 Param_grid_LogReg = {'penalty': ['l1','l2'], 'C': np.linspace(0.1,1.1,10)}
13
14
15 summary = summary.append({'Model' : model_name, 'Training Accuracy' : Mdl_train_sc, 'Te
16                             'Test Recall Score' : Score_Recall, 'Test Precision Score' : Sco
17                             ignore_index=True)

```

```

1 # KNN Model
2
3 Mdl = KNeighborsClassifier()
4
5 model_name = "KNN"
6 Mdl_train_sc = "81.10749185667753"

```

```
7 Score_Accuracy = "74.67532467532467"
8 Score_Recall = "57.89473684210527"
9 Score_Precision = "68.7500"
10 roc_auc = "0.71215409865816603"
11
12 Param_grid_kNeigh = {'n_neighbors': list(np.arange(3,8)), 'metric': ['euclidean','manh
13
14 summary = summary.append({'Model' : model_name, 'Training Accuracy' : Mdl_train_sc, 'Te
15                             'Test Recall Score' : Score_Recall, 'Test Precision Score' : Sco
16                             ignore_index=True})

1 #Naive Bayes Model
2
3 Mdl = GaussianNB()
4
5 model_name = "Naive Byes"
6 Mdl_train_sc = "74.2671009771987"
7 Score_Accuracy = "74.025974025977402"
8 Score_Recall = "63.1578947368421"
9 Score_Precision = "65.4545"
10 roc_auc = "0.7178513293543136"
11
12 summary = summary.append({'Model' : model_name, 'Training Accuracy' : Mdl_train_sc, 'Te
13                             'Test Recall Score' : Score_Recall, 'Test Precision Score' : Sco
14                             ignore_index=True})

1 #SVM Model
2
3 Mdl = SVC(probability=True)
4
5 model_name = "SVM"
6 Mdl_train_sc = "81.92182410423453"
7 Score_Accuracy = "83.11688311688312"
8 Score_Recall = "64.9122807017543"
9 Score_Precision = "86.0465"
10 roc_auc = "0.79386335684572255"
11
12 Param_grid_SVC = {'C': np.linspace(0.1,1.1,10), 'kernel': ['linear','poly','rbf',]}
13
14 summary = summary.append({'Model' : model_name, 'Training Accuracy' : Mdl_train_sc, 'Te
15                             'Test Recall Score' : Score_Recall, 'Test Precision Score' : Sco
16                             ignore_index=True})

1 # Decision Tree Model
2
3 Mdl = DecisionTreeClassifier(random_state=1)
4
5 model_name = "Decision Tree"
6 Mdl_train_sc = "100.0"
7 Score_Accuracy = "75.97402597402598"
8 Score_Recall = "73.68421052631578"
9 Score_Precision = "79.22077922077922"
10 roc_auc = "0.7807921866521975"
```

```

11
12 Param_grid_dt = {'criterion':['gini','entropy'],'max_depth': [3, 4, 5, 6, 7, 8],
13                  'min_impurity_decrease': [0.0001, 0.0003, 0.0005, 0.0007, 0.009]}
14
15
16 summary = summary.append({'Model' : model_name, 'Training Accuracy' : Mdl_train_sc, 'Te
17                          'Test Recall Score' : Score_Recall, 'Test Precision Score' : Sco
18                          ignore_index=True)

1 # Random Forest Model
2
3 Mdl = RandomForestClassifier(random_state=1,n_estimators=100)
4
5 model_name = "Random Forest Classifier"
6 Mdl_train_sc = "100.0"
7 Score_Accuracy = "81.16883116883116"
8 Score_Recall = "66.66666666666666"
9 Score_Precision = "79.1667"
10 roc_auc = "0.781786941580756"
11
12 Param_grid_rf = {'criterion':['gini','entropy'],'max_depth': [3, 4, 5, 6, 7, 8],
13                  'min_impurity_decrease': [0.0001, 0.0003, 0.0005, 0.0007, 0.009]}
14
15
16 summary = summary.append({'Model' : model_name, 'Training Accuracy' : Mdl_train_sc, 'Te
17                          'Test Recall Score' : Score_Recall, 'Test Precision Score' : Sco
18                          ignore_index=True)

1 summary

```

	Model	Training Accuracy	Test Accuracy Score	Test Recall Score	Test
0	LogisticRegression	77.36156351791531	77.27272727272727	57.89473684210527	75.2727
1	KNN	81.10749185667753	74.67532467532467	57.89473684210527	
2	Naive Byes	74.2671009771987	74.025974025977402	63.1578947368421	
3	Naive Byes	74.2671009771987	74.025974025977402	63.1578947368421	
4	SVM	81.92182410423453	83.11688311688312	64.9122807017543	
5	Decision Tree	100.0	75.97402597402598	73.68421052631578	79.2207
6	Random Forest	100.0	81.16883116883116	66.66666666666666	

## Conclusion:

Based on the comparison between the various algorithms used, SVM seems to produce the best results to me.

-By P.Akanksh

IBS(ICFAI BUSSINES SCHOOL),Hyderabad

